

## Predicting the price of Avacados



Some relevant columns in the dataset:

- Date - The date of the observation
- AveragePrice - the average price of a single avocado
- type - conventional or organic
- year - the year
- Region - the city or region of the observation
- Total Volume - Total number of avocados sold
- 4046 - Total number of avocados with PLU 4046 sold
- 4225 - Total number of avocados with PLU 4225 sold
- 4770 - Total number of avocados with PLU 4770 sold

## Importing the libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
sns.set()
import warnings
warnings.filterwarnings("ignore")
```

## Import Dataset

```
In [2]: data = pd.read_csv(r'E:\NIT(Data Science)by (prakash senapati sir (kodi))\Data science all class\August all
```

check the data

In [3]: `data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 18249 entries, 0 to 18248
Data columns (total 14 columns):
#   Column              Non-Null Count  Dtype
---  -
0   Unnamed: 0          18249 non-null  int64
1   Date                18249 non-null  object
2   AveragePrice        18249 non-null  float64
3   Total Volume       18249 non-null  float64
4   4046                18249 non-null  float64
5   4225                18249 non-null  float64
6   4770                18249 non-null  float64
7   Total Bags         18249 non-null  float64
8   Small Bags         18249 non-null  float64
9   Large Bags         18249 non-null  float64
10  XLarge Bags        18249 non-null  float64
11  type                18249 non-null  object
12  year                18249 non-null  int64
13  region              18249 non-null  object
dtypes: float64(9), int64(2), object(3)
memory usage: 1.9+ MB
```


There are 3 categorical features and luckily no missing value. Let's explore the data further.

In [4]: `data.head(3)`

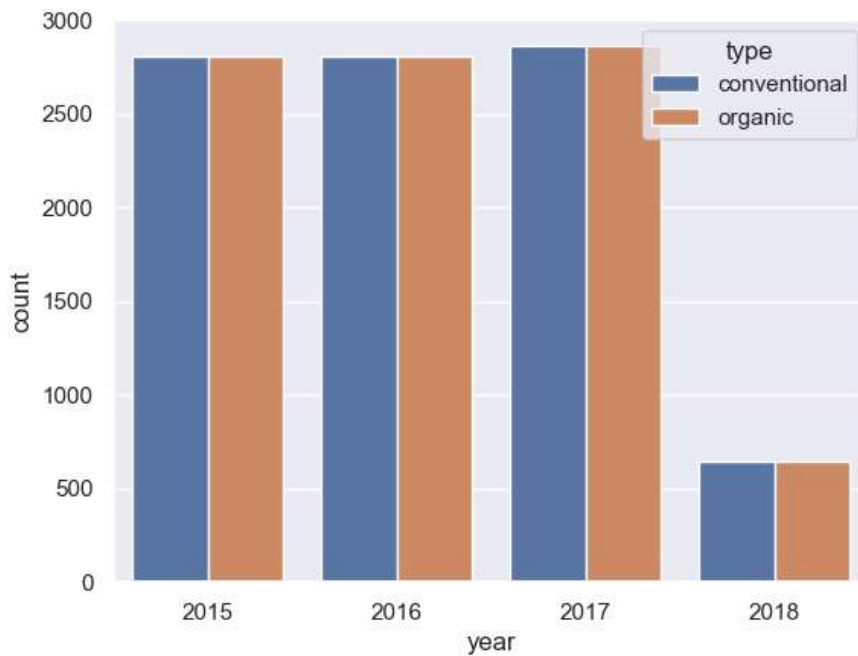
Out[4]:

	Unnamed: 0	Date	AveragePrice	Total Volume	4046	4225	4770	Total Bags	Small Bags	Large Bags	XLarge Bags	type	year	region
0	0	2015-12-27	1.33	64236.62	1036.74	54454.85	48.16	8696.87	8603.62	93.25	0.0	conventional	2015	Albar
1	1	2015-12-20	1.35	54876.98	674.28	44638.81	58.33	9505.56	9408.07	97.49	0.0	conventional	2015	Albar
2	2	2015-12-13	0.93	118220.22	794.70	109149.67	130.50	8145.35	8042.21	103.14	0.0	conventional	2015	Albar

In [5]: `sns.distplot(data["AveragePrice"]);`



```
In [6]: sns.countplot(x="year", data=data, hue="type");
```




There are almost equal numbers of conventional and organic avocados. Though, there is very less observations in the year 2018.

```
In [7]: data.year.value_counts()
```


```
Out[7]: 2017    5722  
        2016    5616  
        2015    5615  
        2018    1296  
        Name: year, dtype: int64
```

```
In [8]: sns.boxplot(y="type", x="AveragePrice", data=data);
```



Organic avocados are more expensive. This is obvious, because their cultivation is more expensive and we all love natural products and are willing to pay a higher price for them.

```
In [9]: data.year=data.year.apply(str)
sns.boxenplot(x="year", y="AveragePrice", data=data);
```



Avacados were slightly more expensive in the year 2017,(as there was shortage due to some reasons)

## Dealing with categorical features

```
In [10]: data["type"] = data["type"].map({"conventional":0, "organic":1})
```

extracting month from data column

```
In [11]: data.Date = data.Date.apply(pd.to_datetime)
data["Month"] = data["Date"].apply(lambda x:x.month)
data.drop("Date", axis=1,inplace=True)
data.Month = data.Month.map({1:"JAN",2:"FEB",3:"MARCH",4:"APRIL",5:"MAY",6:"JUNE",7:"JULY",8:"AUG",9:"SEPT"

```

```
In [12]: ▶ import matplotlib.pyplot as plt
import seaborn as sns

plt.figure(figsize=(9, 5))
sns.barplot(x=data["Month"].value_counts().index, y=data["Month"].value_counts().values)
plt.title("Monthwise Distribution of Sales", fontdict={"fontsize": 25})
plt.xlabel("Month")
plt.ylabel("Number of Sales")
plt.xticks(rotation=45)
plt.show()
```



## Preparing data for ML MODEL

creating dummy variable

```
In [13]: ▶ dummies = pd.get_dummies(data[["year", "region", "Month"]], drop_first=True)
df_dummies = pd.concat([data[["Total Volume", "4046", "4225", "4770", "Total Bags",
"Small Bags", "Large Bags", "XLarge Bags", "type"]], dummies], axis=1)
target = data["AveragePrice"]
```

Splitting data into training and test set

```
In [14]: ▶ from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(df_dummies, target, test_size=0.30)
```

Standardizing the data

```
In [15]: ▶ cols_to_std = ["Total Volume", "4046", "4225", "4770", "Total Bags", "Large Bags", "XLarge Bags"]
from sklearn.preprocessing import StandardScaler
scaler=StandardScaler()
scaler.fit(X_train[cols_to_std])
X_train[cols_to_std] = scaler.transform(X_train[cols_to_std])
X_test[cols_to_std] = scaler.transform(X_test[cols_to_std])
```

## Importing ML models from scikit-learn

```
In [16]: ▶ from sklearn.linear_model import LinearRegression
```

```
In [17]: ▶ from sklearn.tree import DecisionTreeRegressor
```

```
In [18]: ▶ from sklearn.ensemble import RandomForestRegressor
```

```
In [19]: ▶ from sklearn.svm import SVR
```

```
In [20]: ▶ from sklearn.neighbors import KNeighborsRegressor
```

```
In [21]: ▶ pip install xgboost
```

Requirement already satisfied: xgboost in e:\anaconda3\lib\site-packages (1.7.6)  
Requirement already satisfied: numpy in e:\anaconda3\lib\site-packages (from xgboost) (1.23.5)  
Requirement already satisfied: scipy in e:\anaconda3\lib\site-packages (from xgboost) (1.10.0)  
Note: you may need to restart the kernel to use updated packages.

```
In [22]: ▶ from xgboost import XGBRegressor
```

```
In [23]: ▶ from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

to save time all models can be applied once using for loop

```
In [24]: ▶ regressors = {
    "Linear Regression" : LinearRegression(),
    "Decision Tree" : DecisionTreeRegressor(),
    "Random Forest" : RandomForestRegressor(),
    "Support Vector Machines" : SVR(gamma=1),
    "K-nearest Nighbors" : KNeighborsRegressor(n_neighbors=1),
    "XGBoost" : XGBRegressor()
}
results=pd.DataFrame(columns=["MAE", "MSE", "R2-score"])
for method, func in regressors.items():
    model = func.fit(X_train, y_train)
    pred = model.predict(X_test)
    results.loc[method] = [np.round(mean_absolute_error(y_test, pred), 3),
                          np.round(mean_squared_error(y_test, pred), 3),
                          np.round(r2_score(y_test, pred), 3)]
```

## Deep Neural Network

Splitting train set into train and validation sets.

```
In [25]: ▶ X_train, X_val, y_train, y_val = train_test_split(X_train, y_train, test_size=0.20)
```

In [29]: `pip install tensorflow`

```
Requirement already satisfied: tensorflow in e:\anaconda3\lib\site-packages (2.13.0)
Requirement already satisfied: tensorflow-intel==2.13.0 in e:\anaconda3\lib\site-packages (from tensorflow) (2.13.0)
Requirement already satisfied: grpcio<2.0,>=1.24.3 in e:\anaconda3\lib\site-packages (from tensorflow-intel==2.13.0->tensorflow) (1.57.0)
Requirement already satisfied: tensorboard<2.14,>=2.13 in e:\anaconda3\lib\site-packages (from tensorflow-intel==2.13.0->tensorflow) (2.13.0)
Requirement already satisfied: keras<2.14,>=2.13.1 in e:\anaconda3\lib\site-packages (from tensorflow-intel==2.13.0->tensorflow) (2.13.1)
Requirement already satisfied: wrapt>=1.11.0 in e:\anaconda3\lib\site-packages (from tensorflow-intel==2.13.0->tensorflow) (1.14.1)
Requirement already satisfied: opt-einsum>=2.3.2 in e:\anaconda3\lib\site-packages (from tensorflow-intel==2.13.0->tensorflow) (3.3.0)
Requirement already satisfied: protobuf!=4.21.0,!4.21.1,!4.21.2,!4.21.3,!4.21.4,!4.21.5,<5.0.0dev,>=3.20.3 in e:\anaconda3\lib\site-packages (from tensorflow-intel==2.13.0->tensorflow) (4.24.2)
Requirement already satisfied: absl-py>=1.0.0 in e:\anaconda3\lib\site-packages (from tensorflow-intel==2.13.0->tensorflow) (1.4.0)
Requirement already satisfied: h5py>=2.9.0 in e:\anaconda3\lib\site-packages (from tensorflow-intel==2.13.0->tensorflow) (3.7.0)
```

Important Tensorflow libraries

In [27]: `import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Activation, Dropout
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping`

## Creating Model


In [30]: `model = Sequential()
model.add(Dense(76,activation='relu',kernel_initializer=tf.random_uniform_initializer(minval=-0.1, maxval=0.1),
bias_initializer=tf.random_uniform_initializer(minval=-0.1, maxval=0.1)))
model.add(Dense(200,activation='relu',kernel_initializer=tf.random_uniform_initializer(minval=-0.1, maxval=0.1),
bias_initializer=tf.random_uniform_initializer(minval=-0.1, maxval=0.1)))
model.add(Dropout(0.5))
model.add(Dense(200,activation='relu',kernel_initializer=tf.random_uniform_initializer(minval=-0.1, maxval=0.1),
bias_initializer=tf.random_uniform_initializer(minval=-0.1, maxval=0.1)))
model.add(Dropout(0.5))
model.add(Dense(200,activation='relu',kernel_initializer=tf.random_uniform_initializer(minval=-0.1, maxval=0.1),
bias_initializer=tf.random_uniform_initializer(minval=-0.1, maxval=0.1)))
model.add(Dropout(0.5))
model.add(Dense(1))

model.compile(optimizer='Adam', loss='mean_squared_error')
early_stop = EarlyStopping(monitor='val_loss', mode='min', verbose=0, patience=10)`

```
In [31]: model.fit(x=X_train.values,y=y_train.values,
                  validation_data=(X_val.values,y_val.values),
                  batch_size=100,epochs=150,callbacks=[early_stop])
```

Epoch 1/150  
 103/103 [=====] - 5s 13ms/step - loss: 57242268.0000 - val\_loss: 52982.3008  
 Epoch 2/150  
 103/103 [=====] - 1s 9ms/step - loss: 9031492.0000 - val\_loss: 471898.7812  
 Epoch 3/150  
 103/103 [=====] - 1s 10ms/step - loss: 2872031.0000 - val\_loss: 818.7524  
 Epoch 4/150  
 103/103 [=====] - 1s 11ms/step - loss: 911147.7500 - val\_loss: 3975.1584  
 Epoch 5/150  
 103/103 [=====] - 1s 12ms/step - loss: 357832.1875 - val\_loss: 7247.6147  
 Epoch 6/150  
 103/103 [=====] - 1s 11ms/step - loss: 286699.2500 - val\_loss: 3792.0518  
 Epoch 7/150  
 103/103 [=====] - 1s 12ms/step - loss: 76566.8828 - val\_loss: 183.0328  
 Epoch 8/150  
 103/103 [=====] - 1s 12ms/step - loss: 40262.3828 - val\_loss: 37.2701  
 Epoch 9/150  
 103/103 [=====] - 1s 11ms/step - loss: 20512.2773 - val\_loss: 74.5229  
 Epoch 10/150  
 103/103 [=====] - 1s 10ms/step - loss: 5632.3452 - val\_loss: 17.6856

```
In [32]: losses = pd.DataFrame(model.history.history)
losses[["loss","val_loss"]].plot();
```



```
In [33]: dnn_pred = model.predict(X_test)
```

172/172 [=====] - 1s 3ms/step

## Results Table



```
In [35]: results.loc['Deep Neural Network']=[mean_absolute_error(y_test,dnn_pred).round(3),mean_squared_error(y_test,
r2_score(y_test,dnn_pred).round(3)]
results
```

Out[35]:

	MAE	MSE	R2-score
<b>Linear Regression</b>	0.187	0.061	0.628
<b>Decision Tree</b>	0.128	0.037	0.773
<b>Random Forest</b>	0.097	0.019	0.884
<b>Support Vector Machines</b>	0.322	0.161	0.022
<b>K-nearest Nighbors</b>	0.358	0.213	-0.292
<b>XGBoost</b>	0.096	0.017	0.894
<b>Deep Neural Network</b>	0.283	0.131	0.205

```
In [36]: f"10% of mean of target variable is {np.round(0.1 * data.AveragePrice.mean(),3)}"
```

Out[36]: '10% of mean of target variable is 0.141'

```
In [37]: results.sort_values('R2-score',ascending=False).style.background_gradient(cmap='Greens',subset=['R2-score'])
```

Out[37]:

	MAE	MSE	R2-score
<b>XGBoost</b>	0.096000	0.017000	0.894000
<b>Random Forest</b>	0.097000	0.019000	0.884000
<b>Decision Tree</b>	0.128000	0.037000	0.773000
<b>Linear Regression</b>	0.187000	0.061000	0.628000
<b>Deep Neural Network</b>	0.283000	0.131000	0.205000
<b>Support Vector Machines</b>	0.322000	0.161000	0.022000
<b>K-nearest Nighbors</b>	0.358000	0.213000	-0.292000

## Conclusion :-

Except linear regression model, all other models have mean absolute error less than 10% of mean of target variable.

\*For this dataset, XGBoost and Random Forest algorithms have shown best results.

In [ ]: